

Modeling using PCR

We want to see if using principal components help with building the predictive model. The random forest model using original variables is the best we got so far. ## Random Forest Model using Principal Components

```
set.seed(222)
final_previous_merged <- read.csv("final_previous_merged.csv") # load new dataset

final_previous_merged_updated<- final_previous_merged %>% mutate(risk= ifelse(FHR >= 80, "very low risk", "not very low risk"))

# make sure all na values are converted to 0
final_previous_merged_updated[is.na(final_previous_merged_updated)] <- 0

# remove the currency conversion columns
valid_cols <- final_previous_merged_updated %>%
  select(where(is.numeric)) %>%
  select(-prev_X.Other.Currency.to.USD, -prev_inf_factor) %>%
  summarise(across(everything(), ~ mean(!is.na(.)))) %>%
  pivot_longer(everything(), names_to = "col", values_to = "non_na_ratio") %>%
  pull(col)

# also remove FHR to avoid it predicts itself
numeric_data <- final_previous_merged_updated %>%
  select(all_of(valid_cols)) %>%
  select(-FHR) %>%
  drop_na()

nrow(numeric_data)

## [1] 2591

# use Principal Component for modeling, use prcomp function to find number of PCs being used
pca_result <- prcomp(numeric_data, scale. = TRUE)
summary(pca_result)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
## Standard deviation	5.6330	3.1808	1.38322	1.36807	1.27247	1.25719	1.09385
## Proportion of Variance	0.5471	0.1744	0.03299	0.03227	0.02792	0.02725	0.02063
## Cumulative Proportion	0.5471	0.7215	0.75451	0.78678	0.81469	0.84194	0.86257
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
## Standard deviation	1.05610	1.02067	0.99751	0.89910	0.8913	0.74740	0.68511
## Proportion of Variance	0.01923	0.01796	0.01716	0.01394	0.0137	0.00963	0.00809
## Cumulative Proportion	0.88180	0.89977	0.91692	0.93086	0.9446	0.95419	0.96228
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
## Standard deviation	0.65552	0.64473	0.56741	0.43240	0.41759	0.37859	0.34766
## Proportion of Variance	0.00741	0.00717	0.00555	0.00322	0.00301	0.00247	0.00208
## Cumulative Proportion	0.96969	0.97686	0.98241	0.98563	0.98864	0.99111	0.99319
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
## Standard deviation	0.30195	0.25850	0.20352	0.19311	0.17775	0.15506	0.14966
## Proportion of Variance	0.00157	0.00115	0.00071	0.00064	0.00054	0.00041	0.00039
## Cumulative Proportion	0.99476	0.99592	0.99663	0.99727	0.99782	0.99823	0.99862

```
##          PC29    PC30    PC31    PC32    PC33    PC34    PC35
## Standard deviation 0.12137 0.11653 0.09765 0.08924 0.08023 0.07513 0.06934
## Proportion of Variance 0.00025 0.00023 0.00016 0.00014 0.00011 0.00010 0.00008
## Cumulative Proportion 0.99887 0.99911 0.99927 0.99941 0.99952 0.99962 0.99970
##          PC36    PC37    PC38    PC39    PC40    PC41    PC42
## Standard deviation 0.05421 0.05304 0.05013 0.04552 0.04282 0.04216 0.03526
## Proportion of Variance 0.00005 0.00005 0.00004 0.00004 0.00003 0.00003 0.00002
## Cumulative Proportion 0.99975 0.99980 0.99984 0.99988 0.99991 0.99994 0.99996
##          PC43    PC44    PC45    PC46    PC47    PC48    PC49
## Standard deviation 0.02559 0.01968 0.01761 0.0163 0.01298 0.01263 0.01171
## Proportion of Variance 0.00001 0.00001 0.00001 0.0000 0.00000 0.00000 0.00000
## Cumulative Proportion 0.99997 0.99998 0.99998 1.0000 0.99999 0.99999 1.00000
##          PC50    PC51    PC52    PC53    PC54    PC55
## Standard deviation 0.008489 0.006888 0.006081 0.004068 0.003316 0.002259
## Proportion of Variance 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
## Cumulative Proportion 1.000000 1.000000 1.000000 1.000000 1.000000 1.000000
##          PC56    PC57    PC58
## Standard deviation 0.0003631 2.072e-06 3.985e-10
## Proportion of Variance 0.0000000 0.000e+00 0.000e+00
## Cumulative Proportion 1.0000000 1.000e+00 1.000e+00
```

we use 23 Principal Components not the elbow point because we want to maximize the accuracy


```
pca_df <- as.data.frame(pca_result$x[, 1:23])
pca_df$FHR <- final_previous_merged_updated$FHR
set.seed(222)
train_index <- sample(nrow(pca_df), 0.8 * nrow(pca_df))
train <- pca_df[train_index, ]
test <- pca_df[-train_index, ]

rf_model <- randomForest(FHR ~ ., data = train, ntree = 500, importance = TRUE)
rf_pred <- predict(rf_model, newdata = test)
# RMSE
rmse <- sqrt(mean((rf_pred - test$FHR)^2))
cat("RMSE:", round(rmse, 2))
```

```
## RMSE: 9.53
```

```
varImpPlot(rf_model)
```

rf_model



```
top_pcs <- c("PC4", "PC19", "PC18", "PC3")
```

```
actual <- test$FHR
```

```
# R^2
```

```
ss_res <- sum((actual - rf_pred)^2) # residual sum of squares
```

```
ss_tot <- sum((actual - mean(actual))^2) # total sum of squares
```

```
r_squared <- 1 - (ss_res / ss_tot)
```

```
cat("R-squared:", round(r_squared, 4))
```

```
## R-squared: 0.7977
```

```
# this table shows how each top 4 components are consisted by different variables
```

```
pc_loadings <- pca_result$rotation[, top_pcs]
```

```
loading_table <- as.data.frame(pc_loadings) %>%
```

```
  tibble::rownames_to_column("Variable") %>%
```

```
  tidyr::pivot_longer(cols = all_of(top_pcs), names_to = "PC", values_to = "Loading") %>%
```

```
  mutate(abs_loading = abs>Loading)) %>%
```

```
  group_by(PC) %>%
```

```
  slice_max(order_by = abs_loading, n = 10) %>% # top 10 per PC
```


```
  arrange(PC, desc(abs_loading))
```

```
print(loading_table)
```

```
## # A tibble: 40 x 4
```

```
## # Groups:   PC [4]
##   Variable          PC   Loading abs_loading
##   <chr>          <chr>   <dbl>   <dbl>
## 1 prev_debitOwnedWithinOneYear PC18  0.547   0.547
## 2 prev_financialAssets      PC18 -0.464   0.464
## 3 prev_CHS                  PC18  0.365   0.365
## 4 prev_FHR                  PC18 -0.329   0.329
## 5 prev_interestExpense      PC18 -0.271   0.271
## 6 prev_totalCurrentLiabilities PC18  0.209   0.209
## 7 prev_netProfitAfterTax     PC18 -0.192   0.192
## 8 prev_earningsBeforeInterestAndTax PC18 -0.152   0.152
## 9 prev_totalCurrentAssets    PC18  0.136   0.136
## 10 prev_salesRevenue         PC18  0.103   0.103
## # i 30 more rows
```

```
# Show how the prediction fit visually
ggplot(data = NULL, aes(x = test$FHR, y = rf_pred)) +
  geom_point(alpha = 0.6, color = "steelblue") +
  geom_abline(slope = 1, intercept = 0, color = "red", linetype = "dashed") +
  labs(title = "Predicted vs Actual FHR", x = "Actual FHR", y = "Predicted FHR") +
  theme_minimal()
```




First Random Forest model using PCs shows R^2 : 0.7977, RMSE: 9.5322, which shows better than the baseline model, but it is not better than the random forest model using variables.

For the rest of the models, We use the similar code as the first random forest model. ## GLM Model using PCs

```
## RMSE: 9.51
```

```
## R-squared: 0.7988
```




The GLM model using PCs shows R^2 : 0.7988, RMSE: 9.51, which can not beat the random forest model using original variables.

Neural Network Model using PCs

```
## # weights: 126
## initial value 8406633.917903
## iter 10 value 261667.996904
## iter 20 value 231167.170027
## iter 30 value 227040.640039
## iter 40 value 222663.056038
## iter 50 value 221651.813293
## iter 60 value 220766.718154
## iter 70 value 220489.005995
## iter 80 value 220158.788195
## iter 90 value 219461.476704
## iter 100 value 218279.768828
## final value 218279.768828
## stopped after 100 iterations
```

```
## Neural Network RMSE: 9.34
```

R-squared: 0.8059



The Neural Network Model shows R^2 : 0.8059, RMSE: 9.34, which is still not better than the Random Forest Model using the original variables.

In conclusion, after testing different models using PCs, all those PCR models cannot beat the random forest model as predictive models, thus we may want to use Random Forest Model using the original variables to do the next step.