

PCR part

Jiajun Chen

2025-04-22

We want to see if using principal components help with building the predictive model. ## Random Forest Model using Principal Components

```
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2     3.5.0      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(ggplot2)
library(dplyr)
library(randomForest)
```

```
## randomForest 4.7-1.2
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##   combine
##
## The following object is masked from 'package:ggplot2':
##
##   margin
```

```
library(tree)
library(lubridate)
library(grid)
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##
## Attaching package: 'Matrix'
##
## The following objects are masked from 'package:tidyr':
##
##     expand, pack, unpack
##
## Loaded glmnet 4.1-8
```

```
library(corrplot)
```

```
## corrplot 0.95 loaded
```

```
library(tibble)
library(nnet)
set.seed(222)
final_previous_merged <- read.csv("final_previous_merged.csv") # load new dataset

final_previous_merged_updated <- final_previous_merged %>% mutate(risk= ifelse(FHR >= 80, "very low risk", "not very low risk"))

# make sure all na values are converted to 0
final_previous_merged_updated[is.na(final_previous_merged_updated)] <- 0

# remove the currency conversion columns
valid_cols <- final_previous_merged_updated %>%
  select(where(is.numeric)) %>%
  select(-prev_X.Other.Currency.to.USD, -prev_inf_factor) %>%
  summarise(across(everything(), ~ mean(!is.na(.)))) %>%
  pivot_longer(everything(), names_to = "col", values_to = "non_na_ratio") %>%
  pull(col)

# also remove FHR to avoid it predicts itself
numeric_data <- final_previous_merged_updated %>%
  select(all_of(valid_cols)) %>%
  select(-FHR) %>%
  drop_na()

nrow(numeric_data)
```

```
## [1] 2591
```

```
# use Principal Component for modeling, use prcomp function to find number of PCs being used
pca_result <- prcomp(numeric_data, scale. = TRUE)
summary(pca_result)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation    5.6330 3.1808 1.38322 1.36807 1.27247 1.25719 1.09385
## Proportion of Variance 0.5471 0.1744 0.03299 0.03227 0.02792 0.02725 0.02063
## Cumulative Proportion 0.5471 0.7215 0.75451 0.78678 0.81469 0.84194 0.86257
##              PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation    1.05610 1.02067 0.99751 0.89910 0.8913 0.74740 0.68511
```

```
## Proportion of Variance 0.01923 0.01796 0.01716 0.01394 0.0137 0.00963 0.00809
## Cumulative Proportion 0.88180 0.89977 0.91692 0.93086 0.9446 0.95419 0.96228
## PC15 PC16 PC17 PC18 PC19 PC20 PC21
## Standard deviation 0.65552 0.64473 0.56741 0.43240 0.41759 0.37859 0.34766
## Proportion of Variance 0.00741 0.00717 0.00555 0.00322 0.00301 0.00247 0.00208
## Cumulative Proportion 0.96969 0.97686 0.98241 0.98563 0.98864 0.99111 0.99319
## PC22 PC23 PC24 PC25 PC26 PC27 PC28
## Standard deviation 0.30195 0.25850 0.20352 0.19311 0.17775 0.15506 0.14966
## Proportion of Variance 0.00157 0.00115 0.00071 0.00064 0.00054 0.00041 0.00039
## Cumulative Proportion 0.99476 0.99592 0.99663 0.99727 0.99782 0.99823 0.99862
## PC29 PC30 PC31 PC32 PC33 PC34 PC35
## Standard deviation 0.12137 0.11653 0.09765 0.08924 0.08023 0.07513 0.06934
## Proportion of Variance 0.00025 0.00023 0.00016 0.00014 0.00011 0.00010 0.00008
## Cumulative Proportion 0.99887 0.99911 0.99927 0.99941 0.99952 0.99962 0.99970
## PC36 PC37 PC38 PC39 PC40 PC41 PC42
## Standard deviation 0.05421 0.05304 0.05013 0.04552 0.04282 0.04216 0.03526
## Proportion of Variance 0.00005 0.00005 0.00004 0.00004 0.00003 0.00003 0.00002
## Cumulative Proportion 0.99975 0.99980 0.99984 0.99988 0.99991 0.99994 0.99996
## PC43 PC44 PC45 PC46 PC47 PC48 PC49
## Standard deviation 0.02559 0.01968 0.01761 0.0163 0.01298 0.01263 0.01171
## Proportion of Variance 0.00001 0.00001 0.00001 0.0000 0.00000 0.00000 0.00000
## Cumulative Proportion 0.99997 0.99998 0.99998 1.0000 0.99999 0.99999 1.00000
## PC50 PC51 PC52 PC53 PC54 PC55
## Standard deviation 0.008489 0.006888 0.006081 0.004068 0.003316 0.002259
## Proportion of Variance 0.000000 0.000000 0.000000 0.000000 0.000000 0.000000
## Cumulative Proportion 1.000000 1.000000 1.000000 1.000000 1.000000 1.000000
## PC56 PC57 PC58
## Standard deviation 0.0003631 2.072e-06 3.985e-10
## Proportion of Variance 0.0000000 0.000e+00 0.000e+00
## Cumulative Proportion 1.0000000 1.000e+00 1.000e+00
```

```
# we use 23 Principal Components not the elbow point because we want to maximize the accuracy
pca_df <- as.data.frame(pca_result$x[, 1:23])
pca_df$FHR <- final_previous_merged_updated$FHR
set.seed(222)
train_index <- sample(nrow(pca_df), 0.8 * nrow(pca_df))
train <- pca_df[train_index, ]
test <- pca_df[-train_index, ]

rf_model <- randomForest(FHR ~ ., data = train, ntree = 500, importance = TRUE)
rf_pred <- predict(rf_model, newdata = test)
# RMSE
sqrt(mean((rf_pred - test$FHR)^2))
```

```
## [1] 9.532225
```

```
print(rf_model)
```

```
##
## Call:
## randomForest(formula = FHR ~ ., data = train, ntree = 500, importance = TRUE)
## Type of random forest: regression
## Number of trees: 500
```

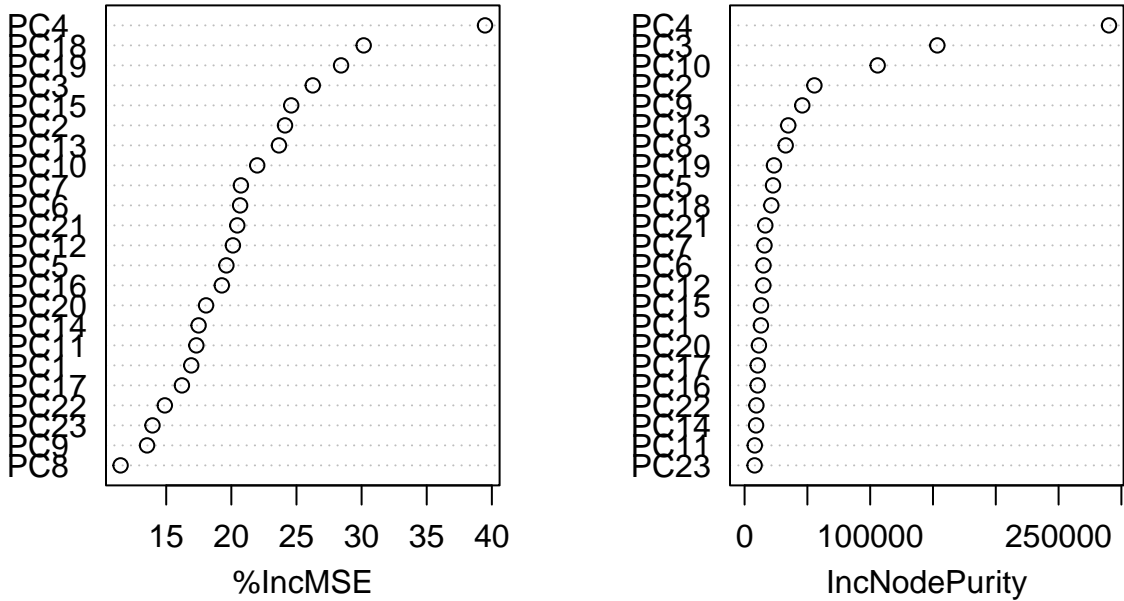
```
## No. of variables tried at each split: 7
##
##           Mean of squared residuals: 88.03068
##           % Var explained: 80.84
```

```
importance(rf_model)
```

```
##           %IncMSE IncNodePurity
## PC1  16.92293    13008.293
## PC2  24.12035    55589.880
## PC3  26.25106   153464.108
## PC4  39.47553   290154.957
## PC5  19.61978    22634.064
## PC6  20.67286    15065.539
## PC7  20.73368    15853.371
## PC8  11.49598    32732.631
## PC9  13.52912    45922.756
## PC10 21.98642   105946.854
## PC11 17.31341     8117.065
## PC12 20.11906    14950.694
## PC13 23.65488    34766.783
## PC14 17.48543     9165.862
## PC15 24.59712    13085.395
## PC16 19.26811    10400.721
## PC17 16.20630    10459.969
## PC18 30.15869    21290.156
## PC19 28.42380    23393.330
## PC20 18.06106    11412.169
## PC21 20.45543    16511.912
## PC22 14.89348     9402.842
## PC23 13.94559     7992.357
```

```
varImpPlot(rf_model)
```

rf_model



```
top_pcs <- c("PC4", "PC19", "PC18", "PC3")
```

```
actual <- test$FHR
```

```
# R^2
```

```
ss_res <- sum((actual - rf_pred)^2) # residual sum of squares
```

```
ss_tot <- sum((actual - mean(actual))^2) # total sum of squares
```

```
r_squared <- 1 - (ss_res / ss_tot)
```

```
cat("R-squared:", round(r_squared, 4))
```

```
## R-squared: 0.7977
```

```
# this table shows how each top 4 components are consisted by different variables
```

```
pc_loadings <- pca_result$rotation[, top_pcs]
```

```
loading_table <- as.data.frame(pc_loadings) %>%
```

```
  tibble::rownames_to_column("Variable") %>%
```

```
  tidyr::pivot_longer(cols = all_of(top_pcs), names_to = "PC", values_to = "Loading") %>%
```

```
  mutate(abs_loading = abs>Loading)) %>%
```

```
  group_by(PC) %>%
```

```
  slice_max(order_by = abs_loading, n = 10) %>% # top 10 per PC
```

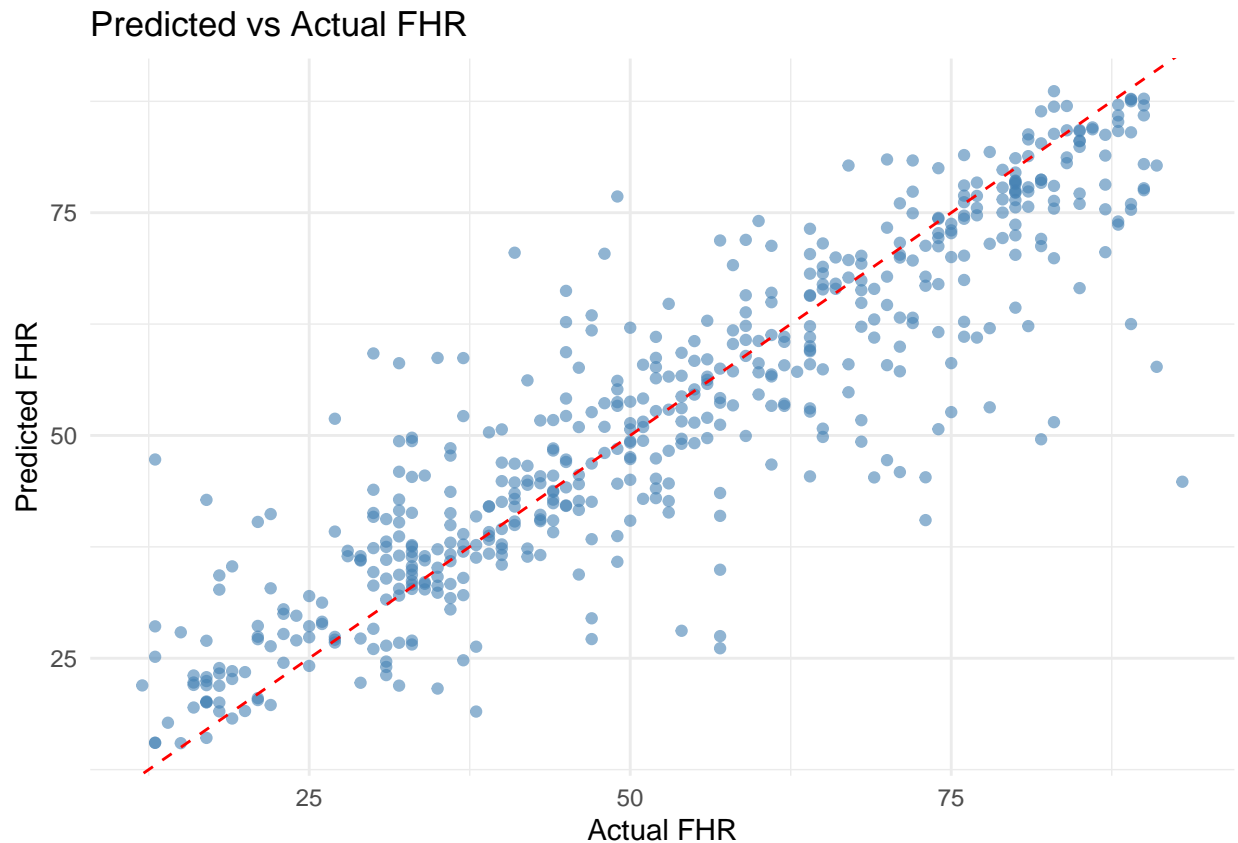
```
  arrange(PC, desc(abs_loading))
```

```
print(loading_table)
```

```
## # A tibble: 40 x 4
```

```
## # Groups:   PC [4]
##   Variable          PC   Loading abs_loading
##   <chr>          <chr>   <dbl>   <dbl>
## 1 prev_debitOwnedWithinOneYear PC18  0.547   0.547
## 2 prev_financialAssets      PC18 -0.464   0.464
## 3 prev_CHS                  PC18  0.365   0.365
## 4 prev_FHR                  PC18 -0.329   0.329
## 5 prev_interestExpense      PC18 -0.271   0.271
## 6 prev_totalCurrentLiabilities PC18  0.209   0.209
## 7 prev_netProfitAfterTax     PC18 -0.192   0.192
## 8 prev_earningsBeforeInterestAndTax PC18 -0.152   0.152
## 9 prev_totalCurrentAssets    PC18  0.136   0.136
## 10 prev_salesRevenue        PC18  0.103   0.103
## # i 30 more rows
```

```
# Show how the prediction fit visually
ggplot(data = NULL, aes(x = test$FHR, y = rf_pred)) +
  geom_point(alpha = 0.6, color = "steelblue") +
  geom_abline(slope = 1, intercept = 0, color = "red", linetype = "dashed") +
  labs(title = "Predicted vs Actual FHR", x = "Actual FHR", y = "Predicted FHR") +
  theme_minimal()
```



First Random Forest model using PCs shows R^2 : 0.7977, RMSE: 9.5322, which shows better than the baseline model.

Random Forest Model using original variables

```
set.seed(222)
train_index_og <- sample(nrow(final_previous_merged), 0.8 * nrow(final_previous_merged))
train_og <- final_previous_merged[train_index, ]
test_og <- final_previous_merged[-train_index, ]

rf_model_og <- randomForest(FHR ~ ., data = train_og, ntree = 500, importance = TRUE)
rf_pred_og <- predict(rf_model_og, newdata = test_og)
sqrt(mean((rf_pred_og - test_og$FHR)^2))
```

```
## [1] 9.25811
```

```
print(rf_model_og)
```

```
##
## Call:
## randomForest(formula = FHR ~ ., data = train_og, ntree = 500,      importance = TRUE)
##               Type of random forest: regression
##               Number of trees: 500
## No. of variables tried at each split: 20
##
##               Mean of squared residuals: 83.19041
##               % Var explained: 81.9
```

```
importance(rf_model_og)
```

	%IncMSE	IncNodePurity
## prev_accountsPayable	9.308327720	3541.2615
## prev_accountsReceivable	9.248558513	2125.5127
## prev_bankCashBalances	13.008445079	7571.2669
## prev_bankOverdraft	-0.005458321	391.7992
## prev_debitOwnedWithinOneYear	18.648597761	13888.5646
## prev_financialAssets	16.084457317	2433.1940
## prev_fixedAssets	16.943727303	5419.8744
## prev_intangibleAssets	5.566473542	1223.3378
## prev_otherCurrentAssets	9.705968567	2242.8208
## prev_otherCurrentLiabilities	11.241606519	3385.1347
## prev_otherEquity	6.818700804	2425.7308
## prev_otherTermAssets	9.831401130	2449.2562
## prev_otherTermLiabilities	6.330539790	2051.6814
## prev_prepayments	6.120930180	1874.5921
## prev_retainedEarnings	10.737086285	8594.9539
## prev_subscribedCapital	7.444793108	2958.0432
## prev_termLoans	14.754298075	6979.9985
## prev_totalAssets	13.271333706	3999.0119
## prev_totalCurrentAssets	13.321986495	4734.6738
## prev_totalCurrentLiabilities	16.717110239	5545.3185
## prev_totalInventories	8.089459727	3435.2306
## prev_totalLiabilities	14.155171978	5930.6321
## prev_totalShareholderEquity	17.417624116	9261.8208

## prev_totalTermAssets	7.731075545	4245.6430
## prev_totalTermLiabilities	8.641410088	5421.4275
## prev_accountPayableAccrualFromCashFlow	2.852316231	2784.2598
## prev_capitalExpenditure	6.651125609	2458.6381
## prev_changeInInventoriesFromCashFlow	6.751559055	3626.7599
## prev_netChangeInCashAndCashEquivalents	3.725022537	1979.6555
## prev_netFundingCashFlow	13.337838982	4031.7932
## prev_netInvestingCashFlow	10.318040797	2414.1040
## prev_unbilledAccountsReceivableRevenueFromCashFlow	6.319060513	3421.0791
## prev_currency	11.584033088	1397.1752
## prev_financialDate	17.570252085	5667.7715
## prev_abnormalItems	5.560176895	2715.1482
## prev_amortization	6.123388842	977.0644
## prev_companyTaxExpense	13.481598864	5852.5482
## prev_costOfGoodsSold	6.055049332	1807.0545
## prev_depreciation	7.646646756	2289.5160
## prev_earningsBeforeInterestAndTax	17.507082981	46202.9883
## prev_grossProfit	7.679841443	2013.3835
## prev_interestExpense	12.348997406	8206.8237
## prev_interestReceived	13.220295984	2257.9607
## prev_investmentIncome	6.831936294	2405.0596
## prev_netInvestmentIncome	10.855509700	9117.2954
## prev_netProfitAfterTax	16.230487442	75311.0532
## prev_netProfitAfterTaxAndMinorityInterests	10.981971650	19460.9272
## prev_netProfitBeforeTax	10.022462703	13959.4770
## prev_netSurplus	12.163367109	28913.1389
## prev_otherIncome	6.157400774	1754.5442
## prev_otherInvestmentExpense	5.791790694	1466.8958
## prev_otherOperatingExpense	9.899617165	2669.6444
## prev_totalOperatingExpense	7.562197221	2866.4843
## prev_totalOperatingRevenue	9.511350301	1962.9969
## prev_totalStaffCosts	8.819463304	2374.6334
## prev_salesRevenue	14.104758425	4292.6880
## prev_eqyYear	12.286862985	3202.4698
## prev_CHS	21.543971116	123661.0969
## prev_FHR	61.906624692	423542.0335
## prev_X.Other.Currency.to.USD	9.532094419	2664.3666
## prev_inf_factor	13.642213455	2973.6192
## diff_days	14.594924249	7507.9784

```

actual_og <- test_og$FHR
ss_res_og <- sum((actual_og - rf_pred_og)^2)
ss_tot_og <- sum((actual_og - mean(actual_og))^2)
r_squared_og <- 1 - (ss_res_og / ss_tot_og)
cat("R-squared:", round(r_squared_og, 4))

```

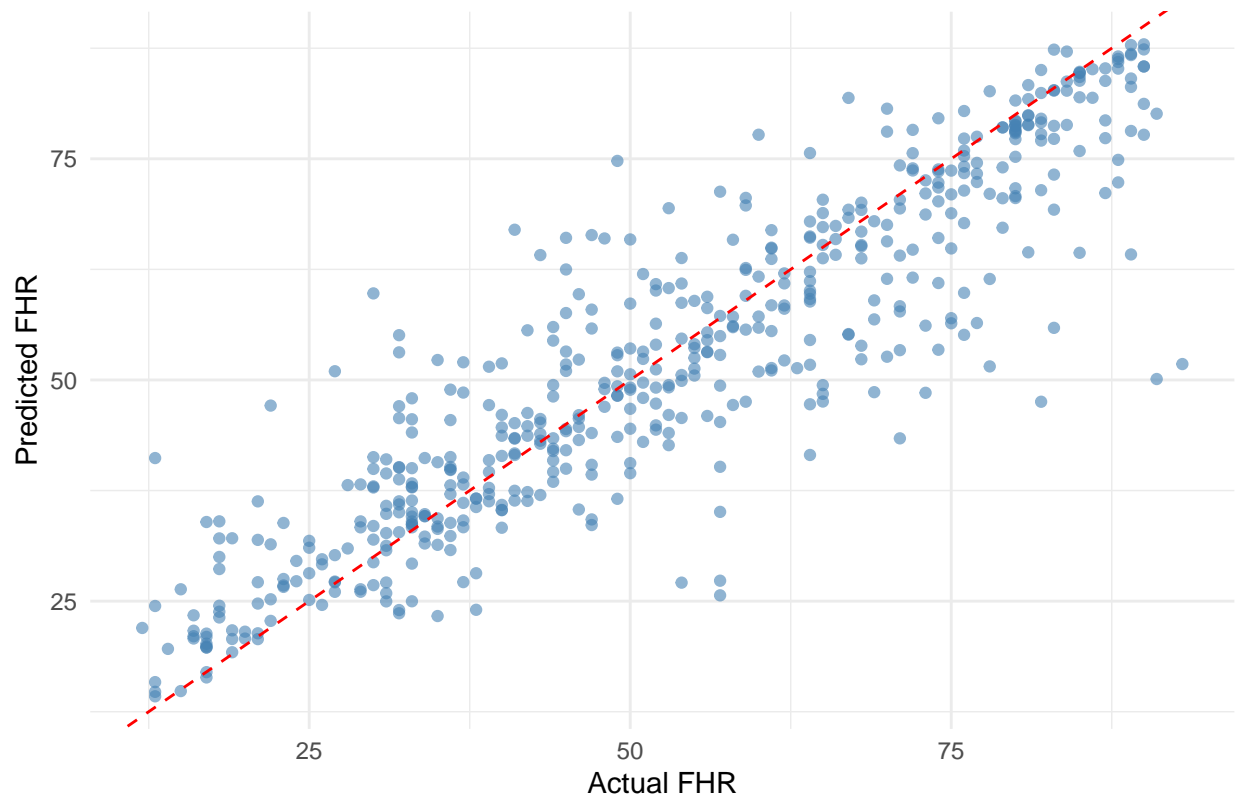
```
## R-squared: 0.8091
```

```

# Show how the prediction fit visually
ggplot(data = NULL, aes(x = test_og$FHR, y = rf_pred_og)) +
  geom_point(alpha = 0.6, color = "steelblue") +
  geom_abline(slope = 1, intercept = 0, color = "red", linetype = "dashed") +
  labs(title = "Predicted vs Actual FHR", x = "Actual FHR", y = "Predicted FHR") +
  theme_minimal()

```


Predicted vs Actual FHR



The Random Forest model using the original variables shows R^2 : 0.8091, RMSE: 9.2581, which shows better than the Random Forest model using PCs, suggests that we may want to stick to this model rather than use random forest model using PCs.

We will try to use principal components for more models to see if it performs better.

GLM Model using PCs

```
set.seed(222)
model_glm <- glm(FHR ~ ., data = train, family = gaussian)
summary(model_glm)
```

```
##
## Call:
## glm(formula = FHR ~ ., family = gaussian, data = train)
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  52.49156    0.20899  251.167  < 2e-16 ***
## PC1          -0.03111    0.03325   -0.936  0.34963
## PC2           1.07768    0.06391  16.862  < 2e-16 ***
## PC3          -4.27523    0.15044 -28.417  < 2e-16 ***
## PC4          11.30062    0.15225  74.222  < 2e-16 ***
## PC5           3.87515    0.16411  23.613  < 2e-16 ***
## PC6          -2.79336    0.16308 -17.129  < 2e-16 ***
```

```
## PC7          -1.55778    0.18303   -8.511   < 2e-16 ***
## PC8           2.15327    0.19056   11.300   < 2e-16 ***
## PC9          -2.73956    0.21258  -12.887   < 2e-16 ***
## PC10         1.31722    0.18967    6.945  5.06e-12 ***
## PC11        -0.43307    0.22863   -1.894   0.05834 .
## PC12        -1.13578    0.22846   -4.971  7.20e-07 ***
## PC13         0.95164    0.44863    2.121   0.03402 *
## PC14        -0.84169    0.32464   -2.593   0.00959 **
## PC15         0.61832    0.31233    1.980   0.04787 *
## PC16        -0.64194    0.31519   -2.037   0.04181 *
## PC17         0.21135    0.40769    0.518   0.60422
## PC18        -6.58001    0.46871  -14.039   < 2e-16 ***
## PC19        -9.12188    0.49431  -18.454   < 2e-16 ***
## PC20         1.05648    0.57249    1.845   0.06512 .
## PC21         4.31424    0.58258    7.405  1.90e-13 ***
## PC22        -0.62442    0.65655   -0.951   0.34169
## PC23         1.20714    0.81562    1.480   0.13902
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 90.29505)
##
##      Null deviance: 952116  on 2071  degrees of freedom
## Residual deviance: 184924  on 2048  degrees of freedom
## AIC: 15236
##
## Number of Fisher Scoring iterations: 2
```

```
glm_preds <- predict(model_glm, newdata = test)
rmse <- sqrt(mean((glm_preds - test$FHR)^2))
cat("RMSE:", round(rmse, 2))
```

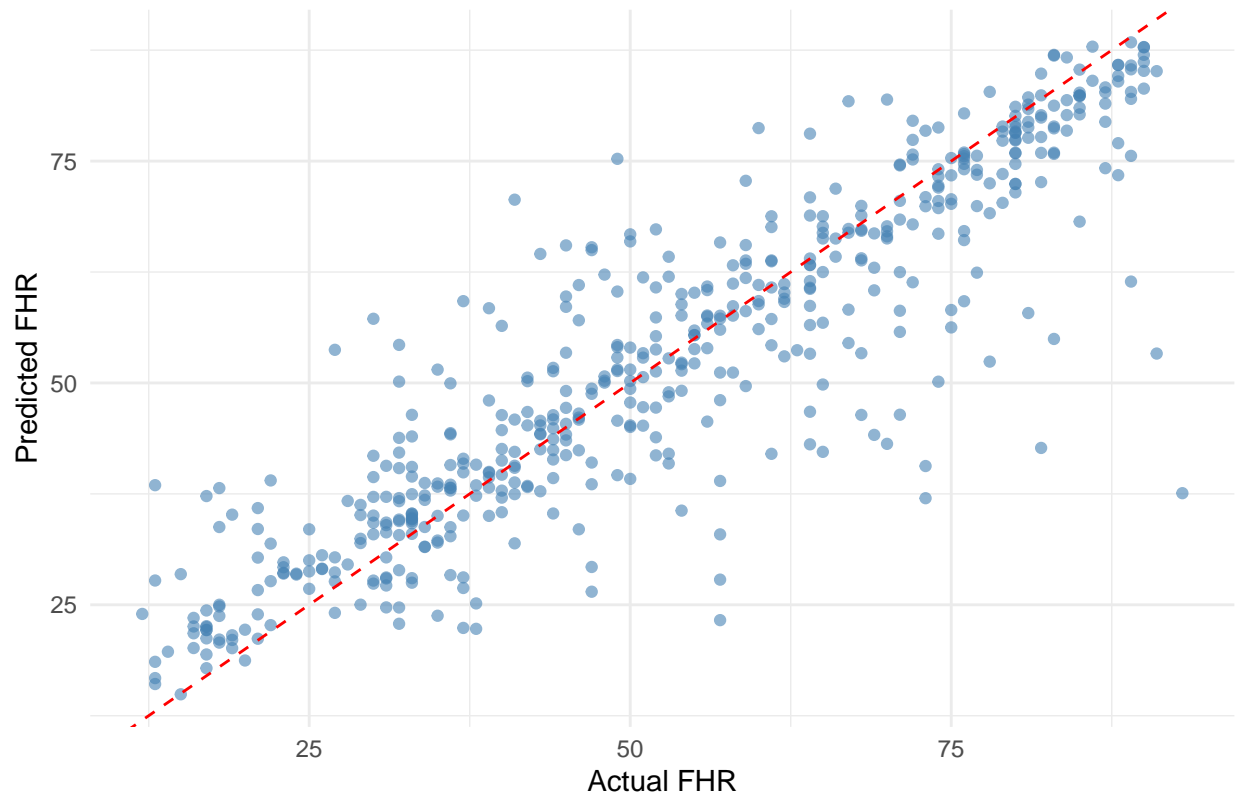
```
## RMSE: 9.51
```

```
actual <- test$FHR
ss_res2 <- sum((actual - glm_preds)^2)
ss_tot2 <- sum((actual - mean(actual))^2)
r_squared2 <- 1 - (ss_res2 / ss_tot2)
cat("R-squared:", round(r_squared2, 4))
```

```
## R-squared: 0.7988
```

```
# Show how the prediction fit visually
ggplot(data = NULL, aes(x = test$FHR, y = glm_preds)) +
  geom_point(alpha = 0.6, color = "steelblue") +
  geom_abline(slope = 1, intercept = 0, color = "red", linetype = "dashed") +
  labs(title = "GLM Predicted vs Actual FHR", x = "Actual FHR", y = "Predicted FHR") +
  theme_minimal()
```

GLM Predicted vs Actual FHR



The GLM model using PCs shows R^2 : 0.7988, RMSE: 9.51, which can not beat the random forest model using original variables.

Neural Network Model using PCs

```
set.seed(222)
nn_model <- nnet(FHR ~ ., data = pca_df, size = 5, linout = TRUE)
```

```
## # weights: 126
## initial value 8406633.917903
## iter 10 value 261667.996904
## iter 20 value 231167.170027
## iter 30 value 227040.640039
## iter 40 value 222663.056038
## iter 50 value 221651.813293
## iter 60 value 220766.718154
## iter 70 value 220489.005995
## iter 80 value 220158.788195
## iter 90 value 219461.476704
## iter 100 value 218279.768828
## final value 218279.768828
## stopped after 100 iterations
```

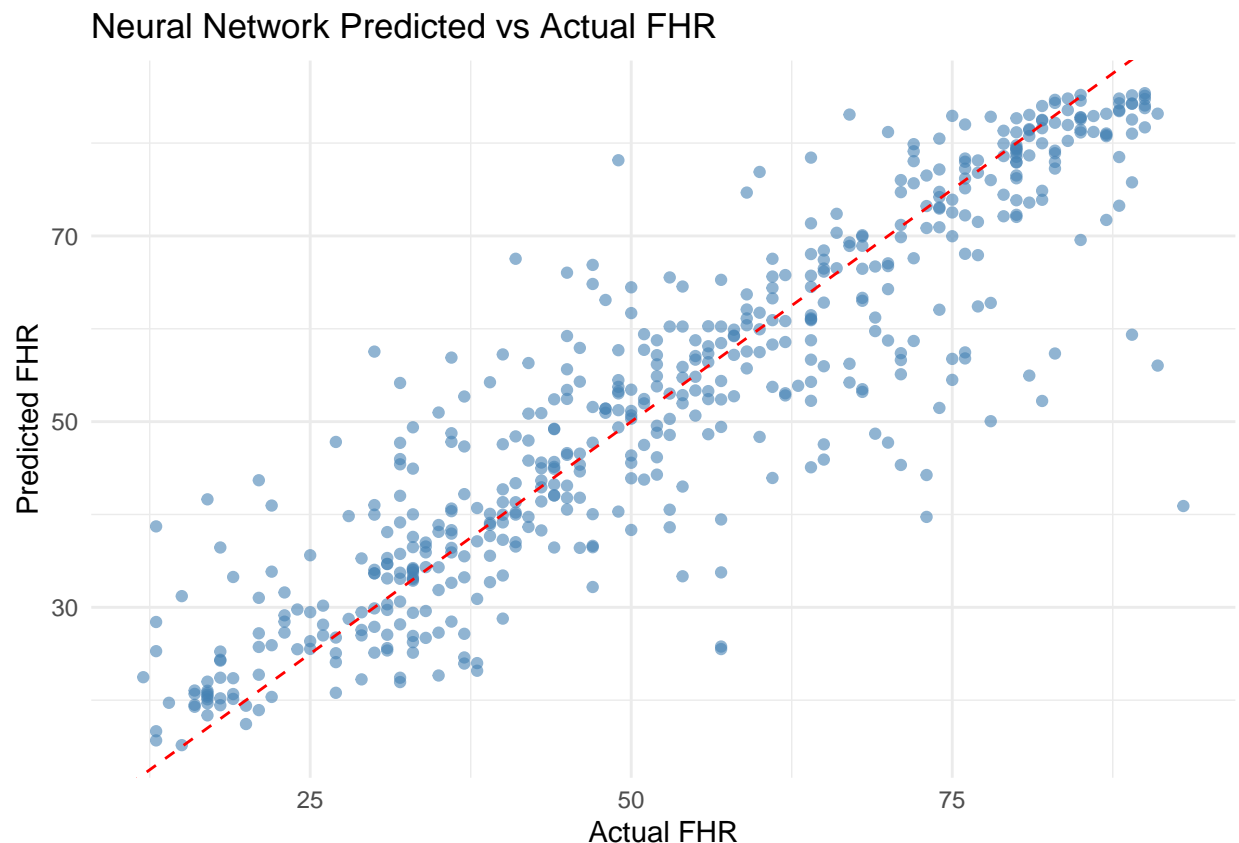
```
nn_preds <- predict(nn_model, newdata = test)
rmse_nn <- sqrt(mean((nn_preds - test$FHR)^2))
cat("Neural Network RMSE:", round(rmse_nn, 2))
```

```
## Neural Network RMSE: 9.34
```

```
actual <- test$FHR
ss_res3 <- sum((actual - nn_preds)^2)
ss_tot3 <- sum((actual - mean(actual))^2)
r_squared2 <- 1 - (ss_res3 / ss_tot3)
cat("R-squared:", round(r_squared2, 4))
```

```
## R-squared: 0.8059
```

```
# Show how the prediction fit visually
ggplot(data = NULL, aes(x = test$FHR, y = nn_preds)) +
  geom_point(alpha = 0.6, color = "steelblue") +
  geom_abline(slope = 1, intercept = 0, color = "red", linetype = "dashed") +
  labs(title = "Neural Network Predicted vs Actual FHR", x = "Actual FHR", y = "Predicted FHR") +
  theme_minimal()
```



The Neural Network Model shows R^2 : 0.8059, RMSE: 9.34, which is still not better than the Random Forest Model using the original variables.

In conclusion, after testing different models using PCs, all those PCR models cannot beat the random forest model as predictive models, thus we may want to use Random Forest Model using the original variables to do the next step.